facebook

facebook

Facebook Architecture

Aditya Agarwal
Director of Engineering
11/22/2008

Agenda

- 1 Architecture Overview
- 2 PHP, MySQL, Memcache
- 3 Thrift, Scribe, Tools
- 4 News Feed Architecture

facebook

At a Glance



The Social Graph

120M+ active users

50B+ PVs per month

10B+ Photos

1B+ connections

50K+ Platform Apps

400K+ App Developers

General Design Principles

- Use open source where possible
 - Explore making optimizations where needed
- Unix Philosophy
 - Keep individual components simple yet performant
 - Combine as necessary
 - Concentrate on clean interface points
- Build everything for scale
- Try to minimize failure points
- Simplicity, Simplicity!

Architecture Overview





LAMP

PHP

Memcache

MySQL

php!

Services

AdServer Search

Network Selector

News Feed

NEWS FEE

Blogfeeds

CSSParser

Mobile

ShareScraper

Thrift Scribe ODS Tools !php

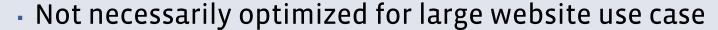
PHP



- Good web programming language
 - Extensive library support for web development
 - Active developer community
- Good for rapid iteration
 - Dynamically typed, interpreted scripting language

PHP: What we Learnt

- Tough to scale for large code bases
 - Weak typing
 - Limited opportunities for static analysis, code optimizations



- E.g. No dynamic reloading of files on web server
- Linearly increasing cost per included file
- Extension framework is difficult to use



PHP: Customizations



- Op-code optimization
- APC improvements
 - Lazy loading
 - Cache priming
 - More efficient locking semantics for variable cache data
- Custom extensions
 - Memcache client extension
 - Serialization format
 - Logging, Stats collection, Monitoring
 - Asynchronous event-handling mechanism

MySQL

Fast, reliable



- Used primarily as <key,value> store
 - Data randomly distributed amongst large set of logical instances
 - Most data access based on global id
- Large number of logical instances spread out across physical nodes
 - Load balancing at physical node level
- No read replication

MySQL: What We Learnt (ing)



Logical migration of data is very difficult

 Create a large number of logical dbs, load balance them over varying number of physical nodes

- No joins in production
 - Logically difficult (because data is distributed randomly)
- Easier to scale CPU on web tier

MySQL: What we Learnt (ing)

My5QL.

- Most data access is for recent data
 - Optimize table layout for recency
 - Archive older data
- Don't ever store non-static data in a central db
 - CDB makes it easier to perform certain aggregated queries
 - Will not scale
- Use services or memcache for global queries
 - E.g.: What are the most popular groups in my network

MySQL: Customizations

No extensive native MySQL modifications



- Custom partitioning scheme
 - Global id assigned to all data
- Custom archiving scheme
 - Based on frequency and recency of data on a per-user basis
- Extended Query Engine for cross-data center replication, cache consistency

MySQL: Customizations



- Graph based data-access libraries
 - Loosely typed objects (nodes) with limited datatypes (int, varchar, text)
 - Replicated connections (edges)
 - Analogous to distributed foreign keys
- Some data collocated
 - Example: User profile data and all of user's connections
- Most data distributed randomly

m e m c a c h e d

Memcache

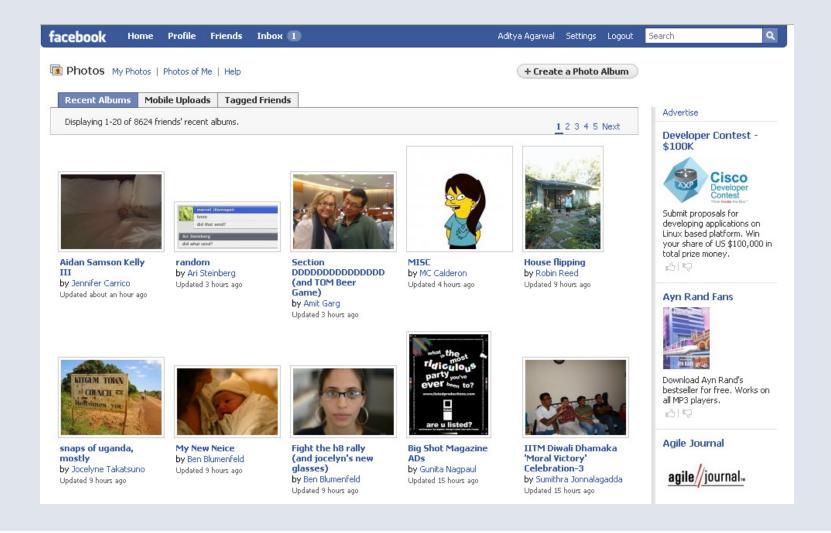
- High-Performance, distributed in-memory hash table
- Used to alleviate database load
- Primary form of caching
- Over 25TB of in-memory cache
- Average latency < 200 micro-seconds
- Cache serialized PHP data structures
- Lots and lots of multi-gets to retrieve data spanning across graph edges

m e m c a c h e d

Memache: Customizations

- Memache over UDP
 - Reduce memory overhead of thousands of TCP connection buffers
 - Application-level flow control (optimization for multi-gets)
- On demand aggregation of per-thread stats
 - Reduces global lock contention
- Multiple Kernel changes to optimize for Memcache usage
 - Distributing network interrupt handling over multiple cores
 - Opportunistic polling of network interface

Let's put this into action



Under the Covers

- Get my profile data
 - Fetch from cache, potentially go to my DB (based on user-id)
- Get friend connections
 - Cache, if not DB (based on user-id)
- In parallel, fetch last 10 photo album ids for each of my friends
 - Multi-get; individual cache misses fetches data from db (based on photoalbum id)
- Fetch data for most recent photo albums in parallel
- Execute page-specific rendering logic in PHP
- Return data, make user happy

LAMP is not Perfect



LAMP is not Perfect

- PHP+MySQL+Memcache works for a large class of problems but not for everything
 - PHP is stateless
 - PHP not the fastest executing language
 - All data is remote
- Reasons why services are written
 - Store code closer to data
 - Compiled environment is more efficient
 - Certain functionality only present in other languages

Services Philosophy

- Create a service iff required
 - Real overhead for deployment, maintenance, separate code-base
 - Another failure point
- Create a common framework and toolset that will allow for easier creation of services
 - Thrift
 - Scribe
 - ODS, Alerting service, Monitoring service
- Use the right language, library and tool for the task

Thrift









High-Level Goal: Enable transparent interaction between these.

...and some others too.



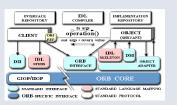
Thrift

- Lightweight software framework for cross-language development
- Provide IDL, statically generate code
- Supported bindings: C++, PHP, Python, Java, Ruby, Erlang, Perl, Haskell etc.
- Transports: Simple Interface to I/O
 - Tsocket, TFileTransport, TMemoryBuffer
- Protocols: Serialization Format
 - TBinaryProtocol, TJSONProtocol
- Servers
 - Non-Blocking, Async, Single Threaded, Multi-threaded

Hasn't this been done before? (yes.)

- SOAP
 - XML, XML, and more XML
- CORBA
 - Bloated? Remote bindings?
- COM
 - Face-Win32ClientSoftware.dll-Book
- Pillar
 - Slick! But no versioning/abstraction.
- Protocol Buffers











Thrift: Why?

- It's quick. Really quick.
- Less time wasted by individual developers
 - No duplicated networking and protocol code
 - Less time dealing with boilerplate stuff
 - Write your client and server in about 5 minutes
- Division of labor
 - Work on high-performance servers separate from applications
- Common toolkit
 - Fosters code reuse and shared tools

Scribe

- Scalable distributed logging framework
- Useful for logging a wide array of data
 - Search Redologs
 - Powers news feed publishing
 - A/B testing data
- Weak Reliability
 - More reliable than traditional logging but not suitable for database transactions.
- Simple data model
- Built on top of Thrift

Other Tools

- SMC (Service Management Console)
 - Centralized configuration
 - Used to determine logical service -> physical node mapping



Other Tools

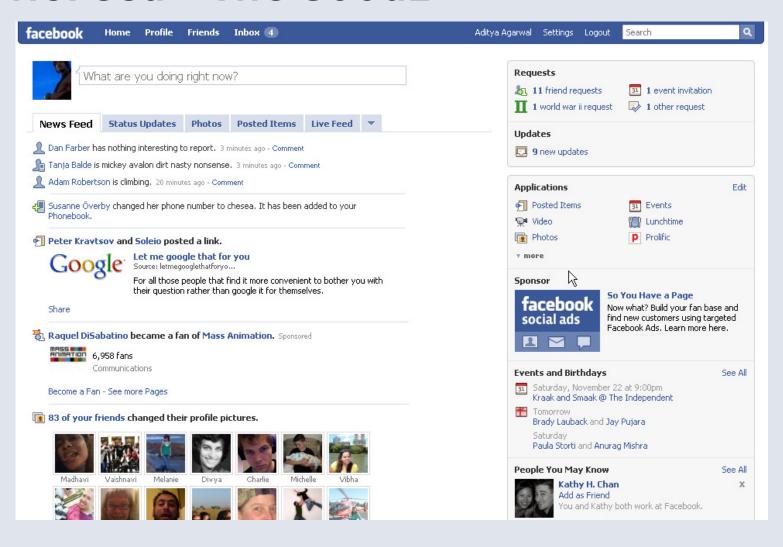
- ODS
 - Used to log and view historical trends for any stats published by service
 - Useful for service monitoring, alerting



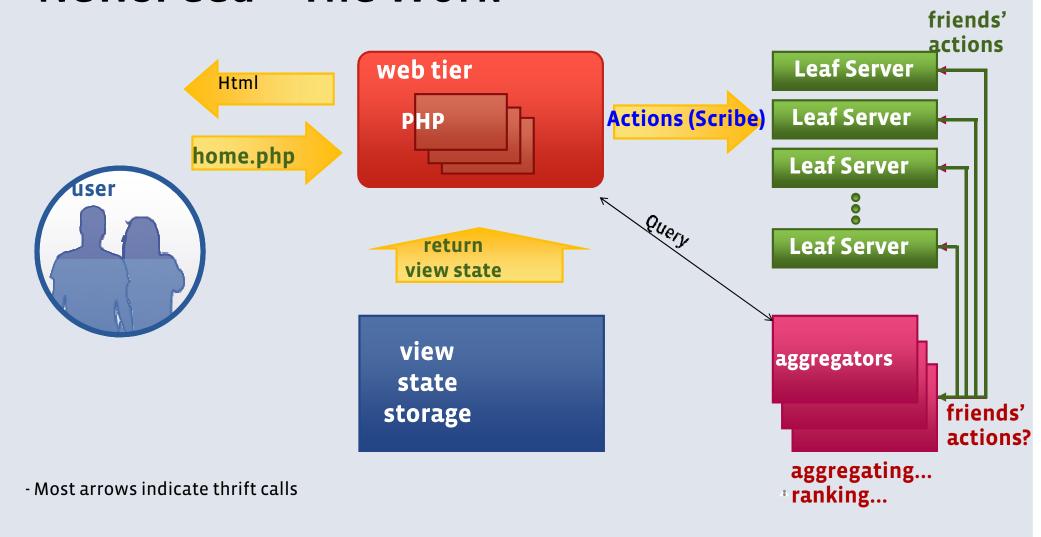
Open Source

- Thrift
 - http://developers.facebook.com/thrift/
- Scribe
 - http://developers.facebook.com/scribe/
- PHPEmbed
 - http://developers.facebook.com/phpembed/
- More good stuff
 - http://developers.facebook.com/opensource.php

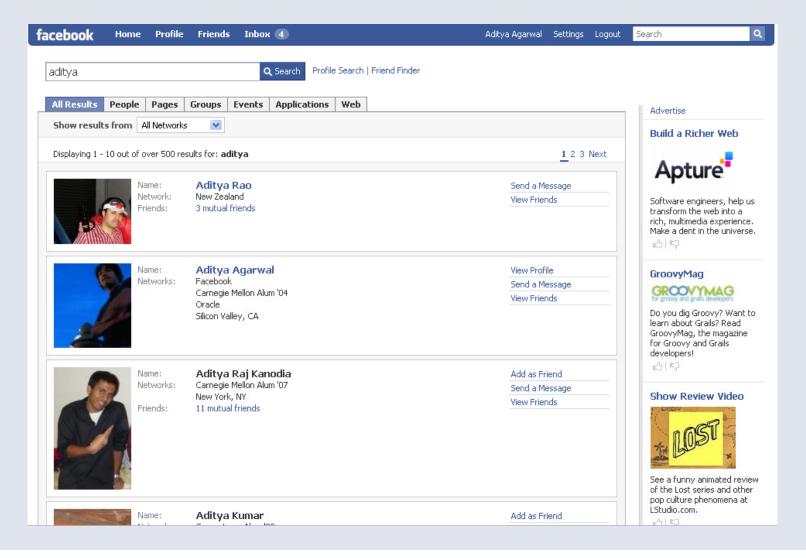
NewsFeed - The Goodz



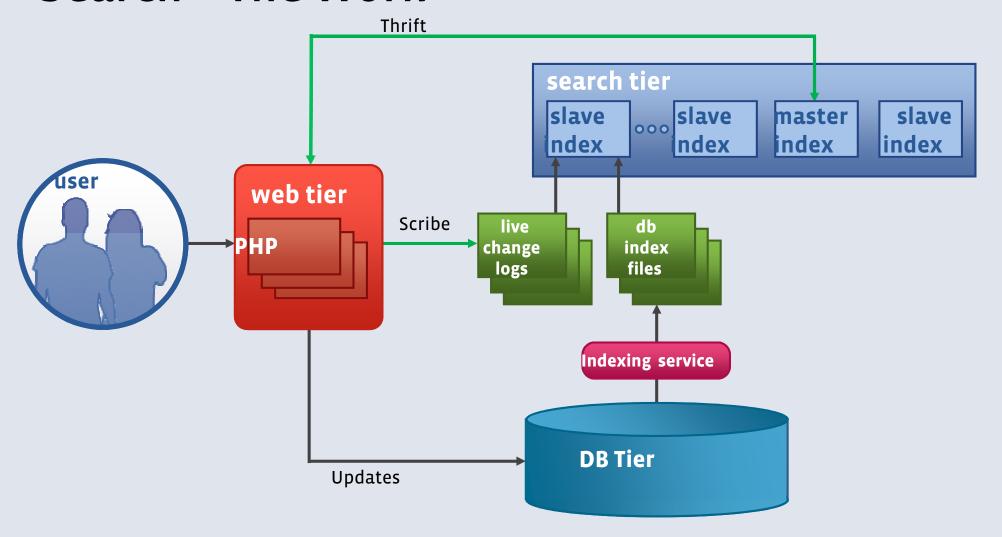
NewsFeed - The Work



Search - The Goodz



Search - The Work



Questions?

More info at www.facebook.com/eblog

Aditya Agarwal aditya@facebook.com