

Distribute Key-Value Store

&Cassandra



Part I 基础介绍



Distribute KV Store 介绍

- KV Store

- Key-Value Store 是一个存在了很久的数据库模式,
value = store.get(key) store.put(key, value) store.delete(key)

- KV Store的特点


- 功能简单
 - 性能高

- Distribute KV Store


- 把KV分布到多个Server上
 - 提供高并发访问和更大的容量
 - 最新的创新使DKV成为云计算的基础组件取得巨大成功

Distribute KV Store Architecture

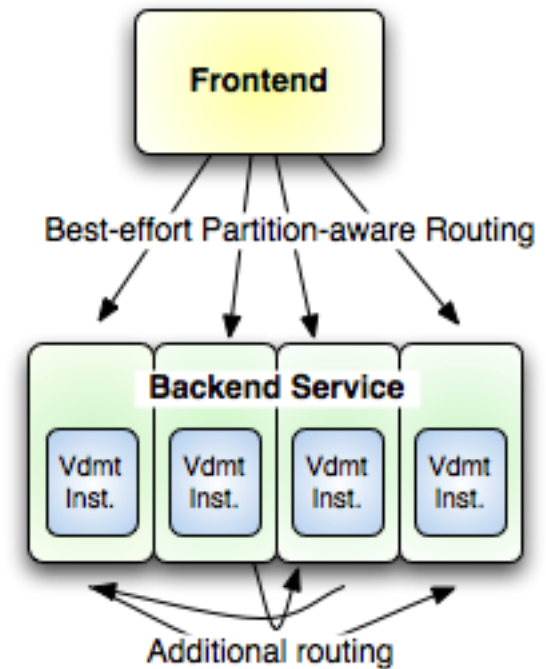
Physical Architecture Options



3-Tier, Server-Routed



3-Tier, Client-Routed



2-Tier, Frontend-Routed

分布式的Key-Value Store和CAP原理

简单的Key-Value一般是嵌入到调用的应用程序中的，所以受限制于机器的内存和磁盘容量，不能服务于海量客户数据。尤其是现在的Web2.0时代，用户创建的数据非常多，不但用Key-Value Store难以存储，即便是使用RDBMS数据库系统，例如Oracle或者Mysql，也难以满足海量数据存放和海量并发读取的要求，或者需要非常高的硬件配置和开二次发代价才能让系统正常工作。

在解决海量数据和海量并发双难题的过程中，业界的一些技术上领先的公司进展很大，象Amazon, Google, Yahoo都分别通过不同的方法解决了问题，他们的方法归结起来就是依靠一点：Distribute, Amazon出了著名的论文：<Dynamo: Amazon's Highly Available Key-Value Store>, 并且在他们的云计算服务EC2/S3中实现了比Dynamo更强大的Simple DB, Google组合了GFS和一系列其他的分布式算法，最终实现了BigTable, Yahoo既有分布式Key-Value Store的实现，也资助了类似于GFS/BigTable的Hadoop项目，到目前为止，Hadoop最大的生产集群依然由Yahoo建立并维护，多达1000 个Nodes.


CAP

在分布式系统中，有一个著名的CAP经验原理，说的是，对于一个分布式系统， consistency(数据一致性)， Availability(可用性)， Partitions(网络分隔)这三者只能同时满足其中的两个。

The CAP Theorem



inktoml



Theorem: You can have **at most two** of these properties for any shared-data system

用CAP理论来解释RDBMS，它满足了Consistency和Availability，但正是因为如此，所以它在Partition上就很难做得好。而对于很多的key-value形式的存储系统而言，它更强调的是Availability和Network Partition，所以在Consistency上做了弱化。例如，Amazon的Dynamo，它就不支持事务，只能提供Eventual Consistency，为了高可用性而牺牲了数据的一致性

“Dynamo targets applications that operate with weaker consistency (the “C” in ACID) if this results in high availability”。

常见的DKV实现

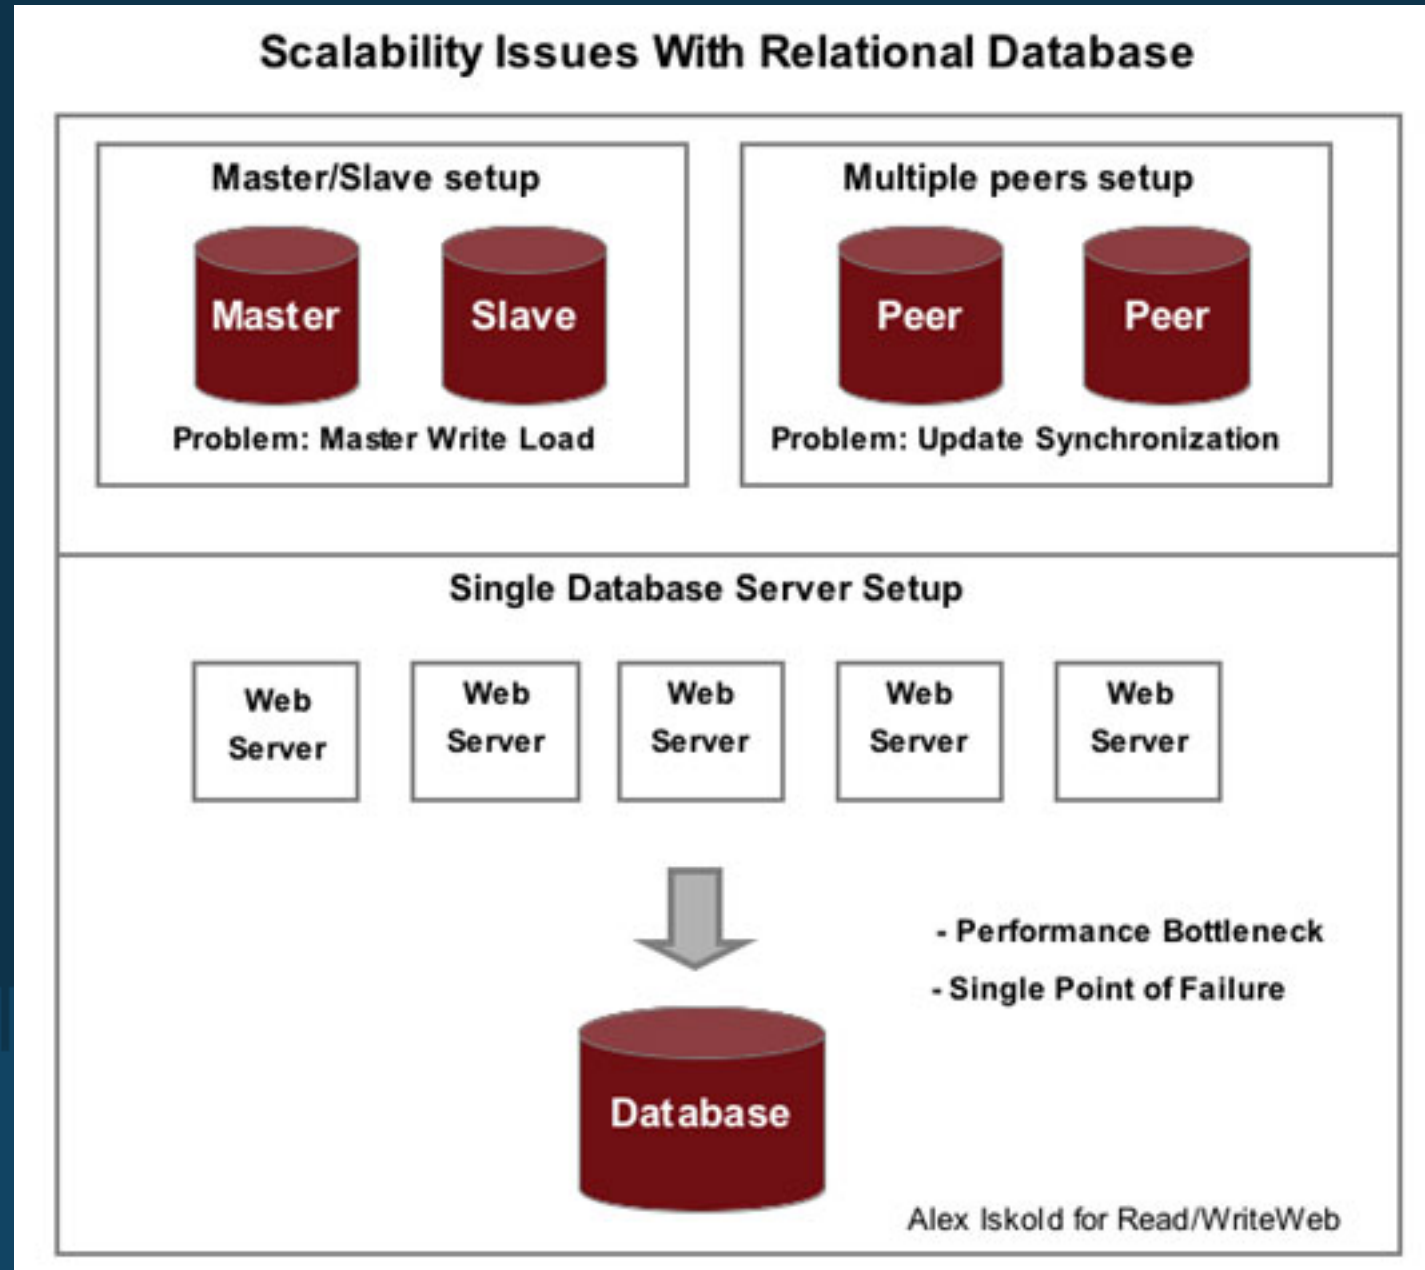
- Amazon Dynamo(下面有详细介绍)
 - Amazon Simple DB
 - Voldmont
 - Cassandra
- Google Big Table
 - HBase
 - Cassandra
 - Google App Engine
- Memcached
 - MemcacheDB

Amazon Dynamo

- Scalability Issues With Relational Databases

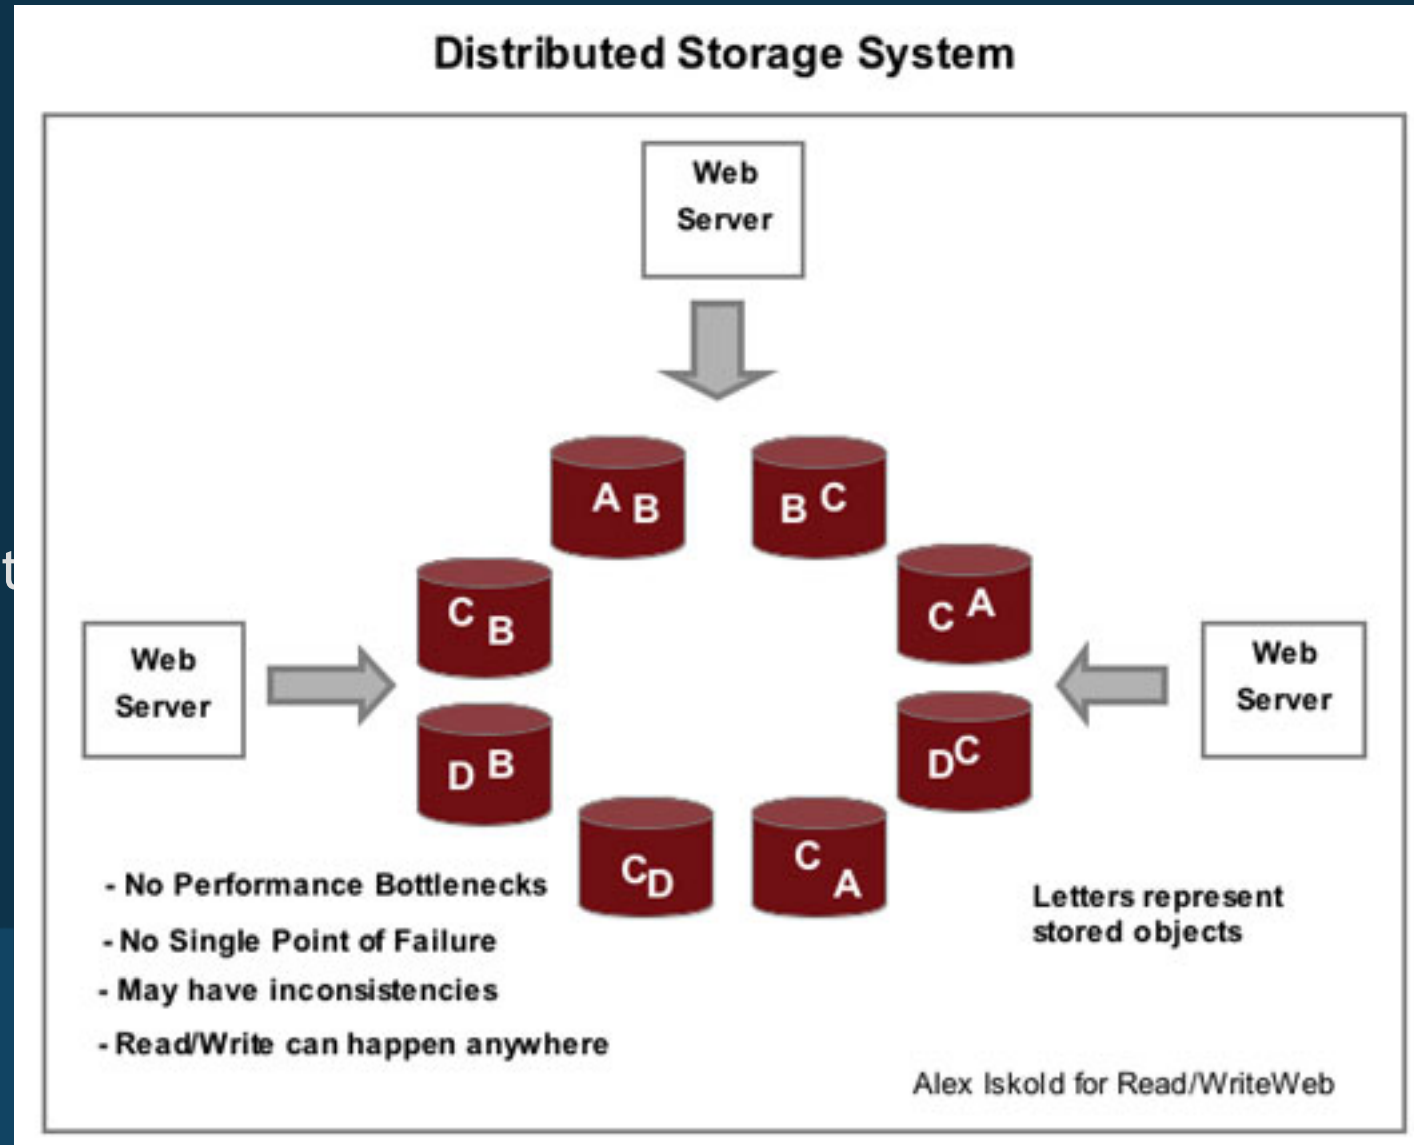
It is difficult to create redundancy and parallelism with relational databases

So as a relational database grows, it becomes a bottleneck and the point of failure for the entire system.




Dynamo - A Distributed Storage System

- Dynamo is a distributed storage system
- the data is made redundant, so each object is stored in the system multiple times.
- Dynamo is called an eventually consistent storage system



How Dynamo Works

- Physical nodes are thought of as identical and organized into a ring.
 - Virtual nodes are created by the system and mapped onto physical nodes, so that hardware can be swapped for maintenance and failure.
 - The partitioning algorithm is one of the most complicated pieces of the system, it specifies which nodes will store a given object.
 - The partitioning mechanism automatically scales as nodes enter and leave the system.
 - Every object is asynchronously replicated to N nodes.
 - The updates to the system occur asynchronously and may result in multiple copies of the object in the system with slightly different states.
 - The discrepancies in the system are reconciled after a period of time, ensuring eventual consistency.
 - Any node in the system can be issued a put or get request for any key.
- 
- A dark blue silhouette of a city skyline with various building shapes of different heights, located at the bottom of the slide.

Part II

Distributed Key-Value Store的 算法基础

Data partitioning and replication

- Data Partitioning:
 - Data needs to be partitioned across a cluster of servers so that no single server needs to hold the complete data set.
- Replication:
 - Put the data into S partitions (one per server) and store copies of a given key K on R servers.

Consistency

- Tolerate the possibility of inconsistency, and resolve inconsistencies at read time.
 - Two-Phase Commit — This is a locking protocol that involves two rounds of co-ordination between machines. It perfectly consistent, but not failure tolerant, and very slow.
 - Paxos-style consensus — This is a protocol for coming to agreement on a value that is more failure tolerant.
 - Read-repair — The first two approaches prevent permanent inconsistency. This approach involves writing all inconsistent versions, and then at read-time detecting the conflict, and resolving the problems. This involves little co-ordination and is completely failure tolerant, but may require additional application logic to resolve conflicts.
- Dynamo use versioning and read-repair.

Versioning

- A vector clock keeps a counter for each writing server, and allows us to calculate when two versions are in conflict, and when one version succeeds or preceeds another.

A vector clock is a list of server:version pairs:

[1:45,2:3,5:55]

The version indicates that the server was the "master" for that number of writes.

A version $v1$ succeeds a version $v2$ if for all i , $v1_i > v2_i$. If neither $v1 > v2$ nor $v1 < v2$, then $v1$ and $v2$ co-occur, and are in conflict.

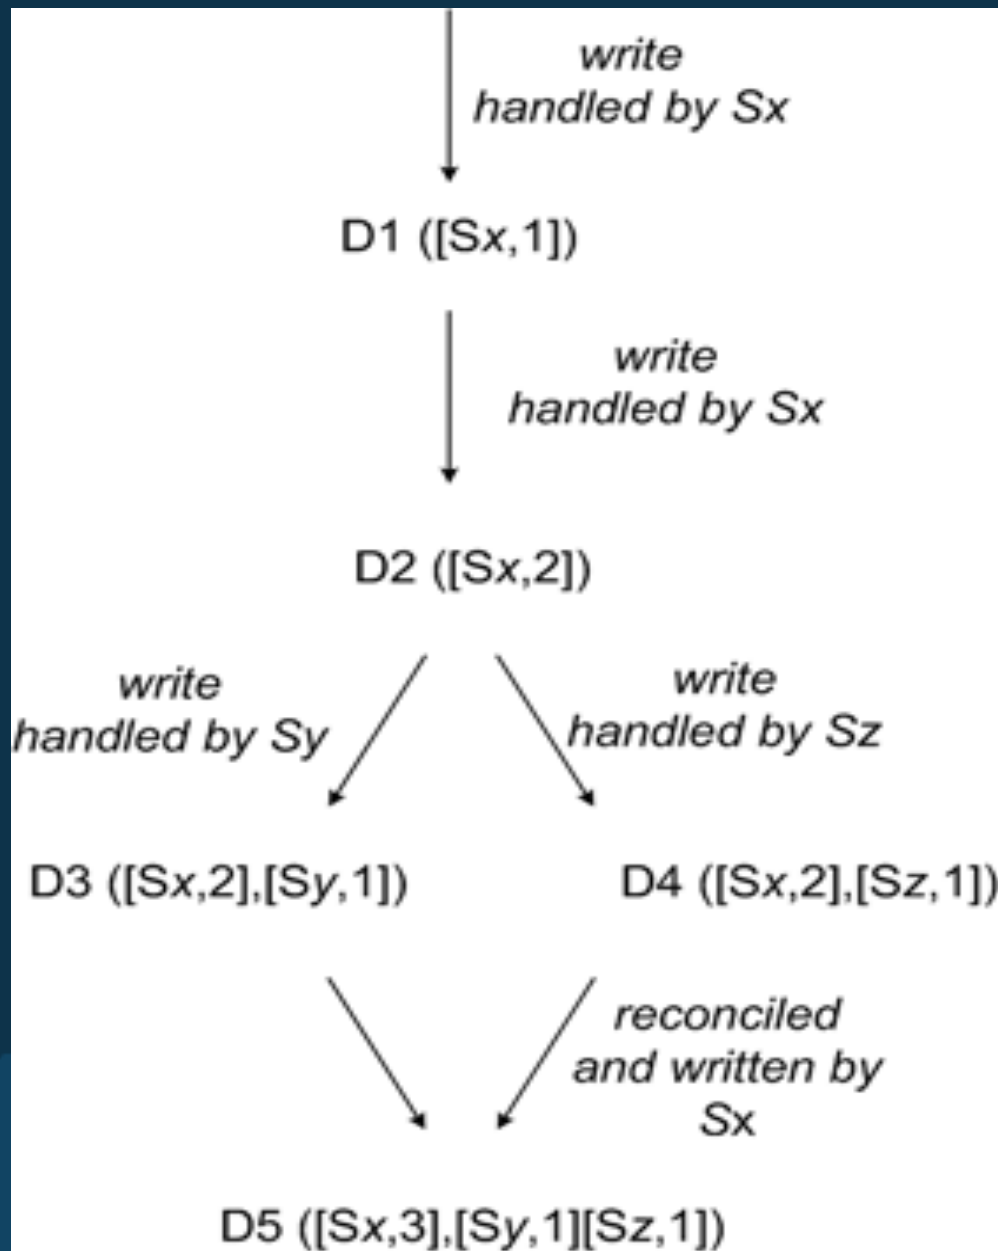
- Here is a simple example of two conflicting versions:

[1:2,2:1]

[1:1,2:2]

- vector clock versioning scheme defines a partial order over values where simple optimistic locking schemes define a total order.

Vector Clock Version



Data partitioning and replication

- Hash Partition


- 使用某种Hash函数分布数据, 性能好, 简单, 缺点在于如果HashBucket发生改变, 会在大范围上导致rehash

- Consistent hashing

- Consistent hashing is a technique that avoids these problems
- When a new server is added to a cluster of S servers, only $1/(S+1)$ values must be moved to the new machine.

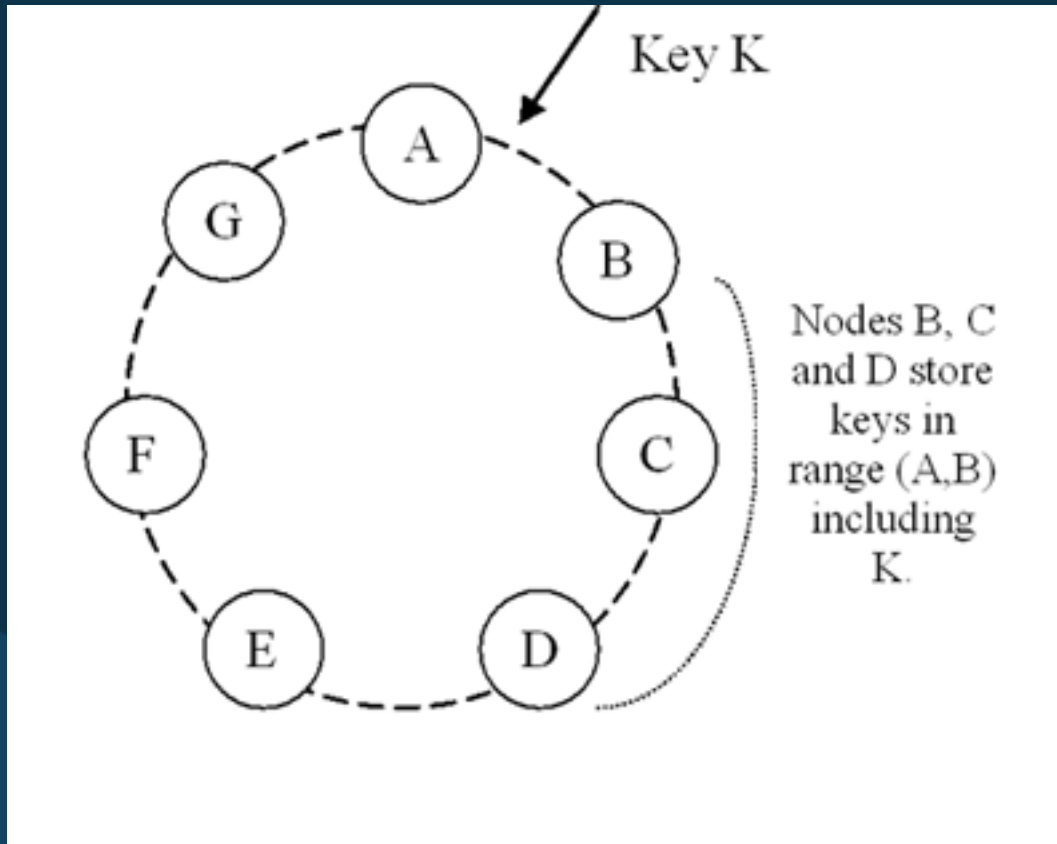
Consistent Hashing

- we can see the possible integer hash values as a ring beginning with 0 and circling around to $2^{31}-1$.
- This ring is divided into Q equally-sized partitions with $Q \gg S$, and each of the S servers is assigned Q/S of these.
- A key is mapped onto the ring using an arbitrary hash function, and then we compute a list of R servers responsible for this key by taking the first R unique nodes when moving over the partitions in a clockwise direction.
- *The diagram below pictures a hash ring for servers A,B,C,D. The arrows indicate keys mapped onto the hash ring and the resulting list of servers that will store the value for that key if $R=3$.*



Replication

To achieve high availability and durability, Dynamo replicates its data on multiple hosts.



Others

- Handling Failures: Hinted Handoff
- Handling permanent failures: Replica synchronization
- HashTree(Merkely Tree)
- ...

<http://s3.amazonaws.com/AllThingsDistributed/sosp/amazon-dynamo-sosp2007.pdf>



Part III DKV的应用



Distributed Key-Value Store的典型应用

- Session Store

SessionID / Session Object

- User注册

UserID / Current Step / Preview Step

- 产品推荐

Dataware Hourse / Greenplum

针对每个用户提供推荐

Cassandra

<http://incubator.apache.org/cassandra>

Cassandra's Features:

- Cassandra is designed to be always available. Writes never fail. Two read paths are available: high-performance "weak" reads and quorum reads. See ThriftInterface.
- Cassandra has a rich data model allowing efficient use for many applications beyond simple key/value.
- Data is automatically replicated to multiple nodes for fault-tolerance. There is support for implementing strategies that replicate across multiple data centers.
- Elasticity: new nodes can be added to a running cluster while minimizing disruption to existing data.
- Consistency: Cassandra follows the "eventually consistent" model but includes sophisticated features such as Hinted Handoff and Read Repair to minimize inconsistency windows.
- Reads and writes in Cassandra are guaranteed to be atomic within a single ColumnFamily.
- Support for versioning and conflict resolution (with inbuilt policies like "last update wins").

Install

- Java1.6
- Apache Thrift
- package 0.4.1
- Only one configure file, very easy to install.

Client Code

```
... TTransport tr = new TSocket("localhost", 9160); TProtocol proto = new TBinaryProtocol(tr);
Cassandra.Client client = new Cassandra.Client(proto); tr.open(); String key_user_id = "1"; // insert
data long timestamp = System.currentTimeMillis(); client.insert("Keyspace1", key_user_id, new
ColumnPath("Standard1", null, "name".getBytes("UTF-8")), "Chris Goffinet".getBytes("UTF-8"), // read
single column ColumnPath path = new ColumnPath("Standard1", null, "name".getBytes("UTF-8"));
System.out.println(client.get("Keyspace1", key_user_id, path, ConsistencyLevel.ONE)); // read entire
row SlicePredicate predicate = new SlicePredicate(null, new SliceRange(new byte[0], new byte[0],
false, 10)); ColumnParent parent = new ColumnParent("Standard1", null); List<ColumnOrSuperColumn>
results = client.get_slice("Keyspace1", key_user_id, parent, predicate, ConsistencyLevel.ONE); for
(ColumnOrSuperColumn result : results) { Column column = result.column; System.out.println(new
String(column.name, "UTF-8") + " -> " + new String(column.value, "UTF-8")); } tr.close(); ...
```

```

service Cassandra {
    # retrieval methods
    ColumnOrSuperColumn get(string keyspace, string key, ColumnPath column_path,
        ConsistencyLevel consistency_level=1)

    list<ColumnOrSuperColumn> get_slice(string keyspace, string key, ColumnParent
        column_parent, SlicePredicate predicate, ConsistencyLevel consistency_level=1)

    map<string, ColumnOrSuperColumn> multiget(string keyspace, list<string> keys, Co
        column_path, ConsistencyLevel consistency_level=1)

    map<string, list<ColumnOrSuperColumn>> multiget_slice(string keyspace,
        list<string> keys, ColumnParent column_parent, SlicePredicate predicate,
        ConsistencyLevel consistency_level=1)

    i32 get_count(string keyspace, string key, ColumnParent column_parent,
        ConsistencyLevel consistency_level=1)

    # range query: returns matching keys
    list<string> get_key_range( string keyspace, string column_family, string
        start="", string finish="", i32 count=100, ConsistencyLevel consistency_level=1)
        throws (1: InvalidRequestException ire, 2: UnavailableException ue),

    # modification methods
    void insert(string keyspace, string key, ColumnPath column_path, binary value,
        i64 timestamp, ConsistencyLevel consistency_level=0)

    void batch_insert(string keyspace string key map<string

```

Cassandra Source Code

- 遵循SEDA模型, Server可以处理大量的并发
- **Memory Table**
- **Serialize all data change action**
- 大量使用异步Callback模型
-

性能测试结果



Thanks

FAQ

